

Neural Networks and Deep Learning Assignment-7

Name :Chanamolu Pranavya

700739974

Github Link: https://github.com/pxc99740/neural_networks_assignment_7.git

Video Link: <https://drive.google.com/drive/folders/1Benn6rUZWV1rPn-B-MmmKgJhRPR1yqpv>

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import RMSprop, Adam
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")
```

```
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
```

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 4s 0us/step
```

```
classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```

```
y_train = y_train.reshape(-1,)
```

```
# Reshape converting 2D to 1D
y_test = y_test.reshape(-1,)
y_train = y_train.reshape(-1,)
```

```
# This code normalization
x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
x_train.shape
```

```
(50000, 32, 32, 3)
```

```
import tensorflow as tf
from tensorflow.keras import layers, models

lenet = models.Sequential([
    layers.Conv2D(6, kernel_size=5, strides=1, activation='relu', input_shape=(32,32,3), padding='same'), #C1
    layers.AveragePooling2D(pool_size=(2, 2)), #S1
    layers.Conv2D(16, kernel_size=5, strides=1, activation='relu', padding='valid'), #C2
    layers.AveragePooling2D(pool_size=(2, 2)), #S2
    layers.Conv2D(120, kernel_size=5, strides=1, activation='relu', padding='valid'), #C3
    layers.Flatten(), #Flatten
    layers.Dense(84, activation='relu'), #F1
    layers.Dense(10, activation='softmax') #Output layer
])
```

```
lenet.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 32, 32, 6)	456
average_pooling2d (Average Pooling2D)	(None, 16, 16, 6)	0
conv2d_1 (Conv2D)	(None, 12, 12, 16)	2416
average_pooling2d_1 (Average Pooling2D)	(None, 6, 6, 16)	0
conv2d_2 (Conv2D)	(None, 2, 2, 120)	48120
flatten (Flatten)	(None, 480)	0
dense (Dense)	(None, 84)	40404
dense_1 (Dense)	(None, 10)	850
=====		
Total params: 92246 (360.34 KB)		
Trainable params: 92246 (360.34 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
lenet.compile(optimizer='adam', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
```

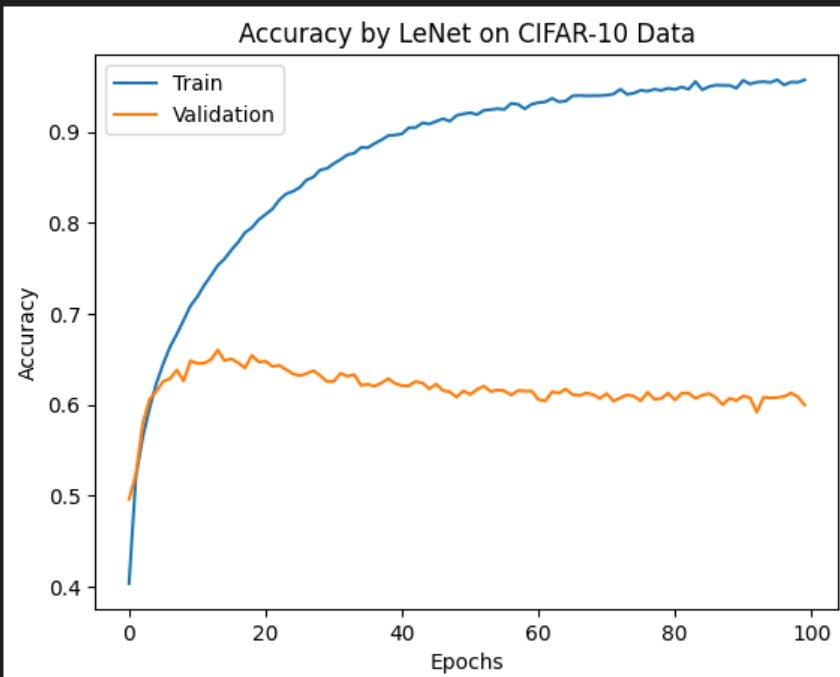
```
hist = lenet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
```


```
Epoch 1/100
1563/1563 [=====] - 14s 6ms/step - loss: 1.6344 - accuracy: 0.4072 - val_loss: 1.3989 - val_accuracy: 0.4972
Epoch 2/100
1563/1563 [=====] - 8s 5ms/step - loss: 1.3263 - accuracy: 0.5271 - val_loss: 1.2305 - val_accuracy: 0.5592
Epoch 3/100
1563/1563 [=====] - 7s 4ms/step - loss: 1.1881 - accuracy: 0.5779 - val_loss: 1.2515 - val_accuracy: 0.5589
Epoch 4/100
1563/1563 [=====] - 8s 5ms/step - loss: 1.1065 - accuracy: 0.6075 - val_loss: 1.1378 - val_accuracy: 0.6011
Epoch 5/100
1563/1563 [=====] - 8s 5ms/step - loss: 1.0338 - accuracy: 0.6352 - val_loss: 1.1170 - val_accuracy: 0.6106
Epoch 6/100
1563/1563 [=====] - 7s 5ms/step - loss: 0.9818 - accuracy: 0.6535 - val_loss: 1.0541 - val_accuracy: 0.6289
Epoch 7/100
1563/1563 [=====] - 9s 5ms/step - loss: 0.9301 - accuracy: 0.6720 - val_loss: 1.0452 - val_accuracy: 0.6337
Epoch 8/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.8856 - accuracy: 0.6885 - val_loss: 1.0623 - val_accuracy: 0.6276
Epoch 9/100
1563/1563 [=====] - 7s 5ms/step - loss: 0.8494 - accuracy: 0.7026 - val_loss: 1.0236 - val_accuracy: 0.6418
Epoch 10/100
1563/1563 [=====] - 7s 5ms/step - loss: 0.8141 - accuracy: 0.7122 - val_loss: 1.0748 - val_accuracy: 0.6269
Epoch 11/100
1563/1563 [=====] - 7s 5ms/step - loss: 0.7773 - accuracy: 0.7258 - val_loss: 1.0560 - val_accuracy: 0.6431
Epoch 12/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.7421 - accuracy: 0.7395 - val_loss: 1.0809 - val_accuracy: 0.6367
Epoch 13/100
...
1563/1563 [=====] - 8s 5ms/step - loss: 0.1506 - accuracy: 0.9472 - val_loss: 3.6584 - val_accuracy: 0.6045
Epoch 71/100
1563/1563 [=====] - 8s 5ms/step - loss: 0.1486 - accuracy: 0.9488 - val_loss: 3.7648 - val_accuracy: 0.6010
Epoch 72/100
1101/1563 [=====>.....] - ETA: 2s - loss: 0.1266 - accuracy: 0.9556
```

```
import numpy as np
```

```
# fix random seed for reproducibility  
seed = 7  
np.random.seed(seed)
```

```
# summarize history for accuracy  
plt.plot(hist.history['accuracy'])  
plt.plot(hist.history['val_accuracy'])  
plt.title("Accuracy by LeNet on CIFAR-10 Data")  
plt.ylabel('Accuracy')  
plt.xlabel('Epochs')  
plt.legend(['Train', 'Validation'], loc='upper left')  
plt.show()  
# summarize history for loss  
plt.plot(hist.history['loss'])  
plt.plot(hist.history['val_loss'])  
plt.title('Loss by LeNet on CIFAR-10 Data')  
plt.ylabel('Loss')  
plt.xlabel('Epochs')  
plt.legend(['Train', 'Validation'])  
plt.show()
```






```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1,)
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)
```

313/313 [=====] - 1s 2ms/step

```
array([[619, 15, 91, 40, 22, 14, 21, 11, 113, 54],
       [ 23, 731, 21, 16, 6, 12, 19, 5, 48, 119],
       [ 57, 6, 531, 97, 86, 87, 71, 30, 21, 14],
       [ 27, 12, 106, 439, 65, 179, 92, 32, 17, 31],
       [ 33, 3, 143, 95, 493, 73, 66, 64, 22, 8],
       [ 13, 8, 131, 211, 67, 453, 41, 50, 11, 15],
       [ 8, 3, 92, 106, 48, 36, 666, 17, 9, 15],
       [ 19, 6, 50, 73, 110, 78, 17, 613, 6, 28],
       [ 80, 30, 30, 27, 19, 6, 21, 4, 748, 35],
       [ 43, 90, 34, 31, 7, 14, 18, 14, 47, 702]])
```

```
# confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```


















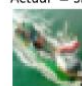





















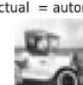












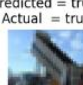
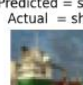



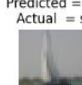
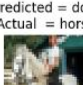

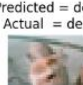
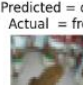
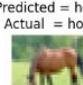
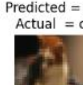


```
print("Test accuracy:", accuracy_score(y_test, y_predictions))
```

Test accuracy: 0.5995

```
L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```

Predicted = dog Actual = cat 	Predicted = ship Actual = ship 	Predicted = airplane Actual = ship 	Predicted = ship Actual = airplane 	Predicted = dog Actual = frog 	Predicted = frog Actual = frog 	Predicted = frog Actual = automobile 	Predicted = frog Actual = frog 
Predicted = cat Actual = cat 	Predicted = truck Actual = automobile 	Predicted = airplane Actual = airplane 	Predicted = truck Actual = truck 	Predicted = dog Actual = dog 	Predicted = horse Actual = horse 	Predicted = truck Actual = truck 	Predicted = airplane Actual = ship 
Predicted = dog Actual = dog 	Predicted = horse Actual = horse 	Predicted = ship Actual = ship 	Predicted = frog Actual = frog 	Predicted = horse Actual = horse 	Predicted = bird Actual = airplane 	Predicted = bird Actual = deer 	Predicted = truck Actual = truck 
Predicted = bird Actual = dog 	Predicted = dog Actual = bird 	Predicted = frog Actual = deer 	Predicted = airplane Actual = airplane 	Predicted = truck Actual = truck 	Predicted = frog Actual = frog 	Predicted = frog Actual = frog 	Predicted = deer Actual = dog 
Predicted = bird Actual = deer 	Predicted = deer Actual = dog 	Predicted = truck Actual = truck 	Predicted = bird Actual = bird 	Predicted = horse Actual = deer 	Predicted = automobile Actual = automobile 	Predicted = truck Actual = truck 	Predicted = dog Actual = dog 
Predicted = deer Actual = deer 	Predicted = frog Actual = frog 	Predicted = dog Actual = dog 	Predicted = frog Actual = frog 	Predicted = airplane Actual = airplane 	Predicted = truck Actual = truck 	Predicted = cat Actual = cat 	Predicted = airplane Actual = truck 
Predicted = horse Actual = horse 	Predicted = frog Actual = frog 	Predicted = truck Actual = truck 	Predicted = ship Actual = ship 	Predicted = frog Actual = airplane 	Predicted = cat Actual = cat 	Predicted = ship Actual = ship 	Predicted = ship Actual = ship 
Predicted = dog Actual = horse 	Predicted = horse Actual = horse 	Predicted = deer Actual = deer 	Predicted = cat Actual = frog 	Predicted = horse Actual = horse 	Predicted = cat Actual = cat 	Predicted = frog Actual = frog 	Predicted = cat Actual = cat 

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, Dropout, Flatten
from tensorflow.keras.optimizers import SGD, Adam
from tensorflow.keras.layers import Convolution2D as Conv2D
from tensorflow.keras.layers import MaxPooling2D

#Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(3,3),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(200, activation='relu'))
AlexNet.add(Dropout(0.1))
AlexNet.add(Dense(10,activation='softmax'))

AlexNet.compile(optimizer='SGD', loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 8, 8, 16)	448
max_pooling2d (MaxPooling2D)	(None, 4, 4, 16)	0
conv2d_4 (Conv2D)	(None, 4, 4, 60)	24060
max_pooling2d_1 (MaxPooling2D)	(None, 2, 2, 60)	0
conv2d_5 (Conv2D)	(None, 2, 2, 60)	32460
conv2d_6 (Conv2D)	(None, 2, 2, 30)	16230
conv2d_7 (Conv2D)	(None, 2, 2, 20)	5420
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 20)	0
flatten_1 (Flatten)	(None, 20)	0
...		
Total params: 125028 (488.39 KB)		
Trainable params: 125028 (488.39 KB)		
Non-trainable params: 0 (0.00 Byte)		


```
history1 = AlexNet.fit(x_train, y_train, epochs=100, validation_data=(x_test, y_test), verbose=1)
```

```
Epoch 1/100
1563/1563 [=====] - 13s 6ms/step - loss: 2.3005 - accuracy: 0.1202 - val_loss: 2.2954 - val_accuracy: 0.1593
Epoch 2/100
1563/1563 [=====] - 9s 6ms/step - loss: 2.2379 - accuracy: 0.1729 - val_loss: 2.0964 - val_accuracy: 0.2019
Epoch 3/100
1563/1563 [=====] - 9s 6ms/step - loss: 2.0718 - accuracy: 0.2052 - val_loss: 2.0534 - val_accuracy: 0.2278
Epoch 4/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.9807 - accuracy: 0.2478 - val_loss: 2.0285 - val_accuracy: 0.2387
Epoch 5/100
1563/1563 [=====] - 10s 6ms/step - loss: 1.8673 - accuracy: 0.2863 - val_loss: 1.7604 - val_accuracy: 0.3378
Epoch 6/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.7570 - accuracy: 0.3271 - val_loss: 1.7866 - val_accuracy: 0.3321
Epoch 7/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.6764 - accuracy: 0.3639 - val_loss: 1.6310 - val_accuracy: 0.3875
Epoch 8/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.6125 - accuracy: 0.3900 - val_loss: 1.6049 - val_accuracy: 0.3961
Epoch 9/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.5451 - accuracy: 0.4206 - val_loss: 1.5705 - val_accuracy: 0.4325
Epoch 10/100
1563/1563 [=====] - 9s 5ms/step - loss: 1.4799 - accuracy: 0.4498 - val_loss: 1.6764 - val_accuracy: 0.3789
Epoch 11/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.4215 - accuracy: 0.4730 - val_loss: 1.3958 - val_accuracy: 0.4890
Epoch 12/100
1563/1563 [=====] - 9s 6ms/step - loss: 1.3713 - accuracy: 0.4935 - val_loss: 1.3535 - val_accuracy: 0.5073
Epoch 13/100
...
Epoch 99/100
1563/1563 [=====] - 9s 6ms/step - loss: 0.2739 - accuracy: 0.9039 - val_loss: 2.5997 - val_accuracy: 0.5227
Epoch 100/100
1563/1563 [=====] - 9s 6ms/step - loss: 0.2681 - accuracy: 0.9057 - val_loss: 2.5864 - val_accuracy: 0.5327
```

```
# summarize history for accuracy
plt.plot(history1.history['accuracy'])
plt.plot(history1.history['val_accuracy'])
plt.title("Accuracy by AlexNet on CIFAR-10 Data")
plt.ylabel('Accuracy')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# summarize history for loss
plt.plot(history1.history['loss'])
plt.plot(history1.history['val_loss'])
plt.title('Loss by AlexNet on CIFAR-10 Data')
plt.ylabel('Loss')
plt.xlabel('Epochs')
plt.legend(['Train', 'Validation'])
plt.show()
```


```
y_predictions1 = AlexNet.predict(x_test)
y_predictions1.reshape(-1,)
y_predictions1= np.argmax(y_predictions1, axis=1)

confusion_matrix(y_test, y_predictions1)
```

313/313 [=====] - 1s 2ms/step

```
array([[636, 32, 69, 48, 32, 13, 9, 41, 66, 54],
       [ 57, 585, 10, 28, 8, 11, 15, 31, 50, 205],
       [ 90, 4, 381, 106, 169, 71, 70, 66, 21, 22],
       [ 33, 16, 73, 449, 75, 202, 70, 53, 10, 19],
       [ 32, 8, 93, 83, 519, 59, 82, 105, 12, 7],
       [ 14, 6, 64, 257, 75, 429, 40, 95, 4, 16],
       [ 7, 12, 78, 136, 146, 53, 515, 16, 11, 26],
       [ 31, 14, 54, 65, 107, 82, 17, 582, 6, 42],
       [147, 52, 26, 42, 30, 15, 13, 9, 617, 49],
       [ 65, 124, 24, 37, 9, 16, 12, 62, 37, 614]])
```

```
# confusion matrix and accuracy
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix of AlexNet Model', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions1))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
print("Test accuracy by AlexNet:", accuracy_score(y_test, y_predictions))
```

Test accuracy by AlexNet: 0.5995

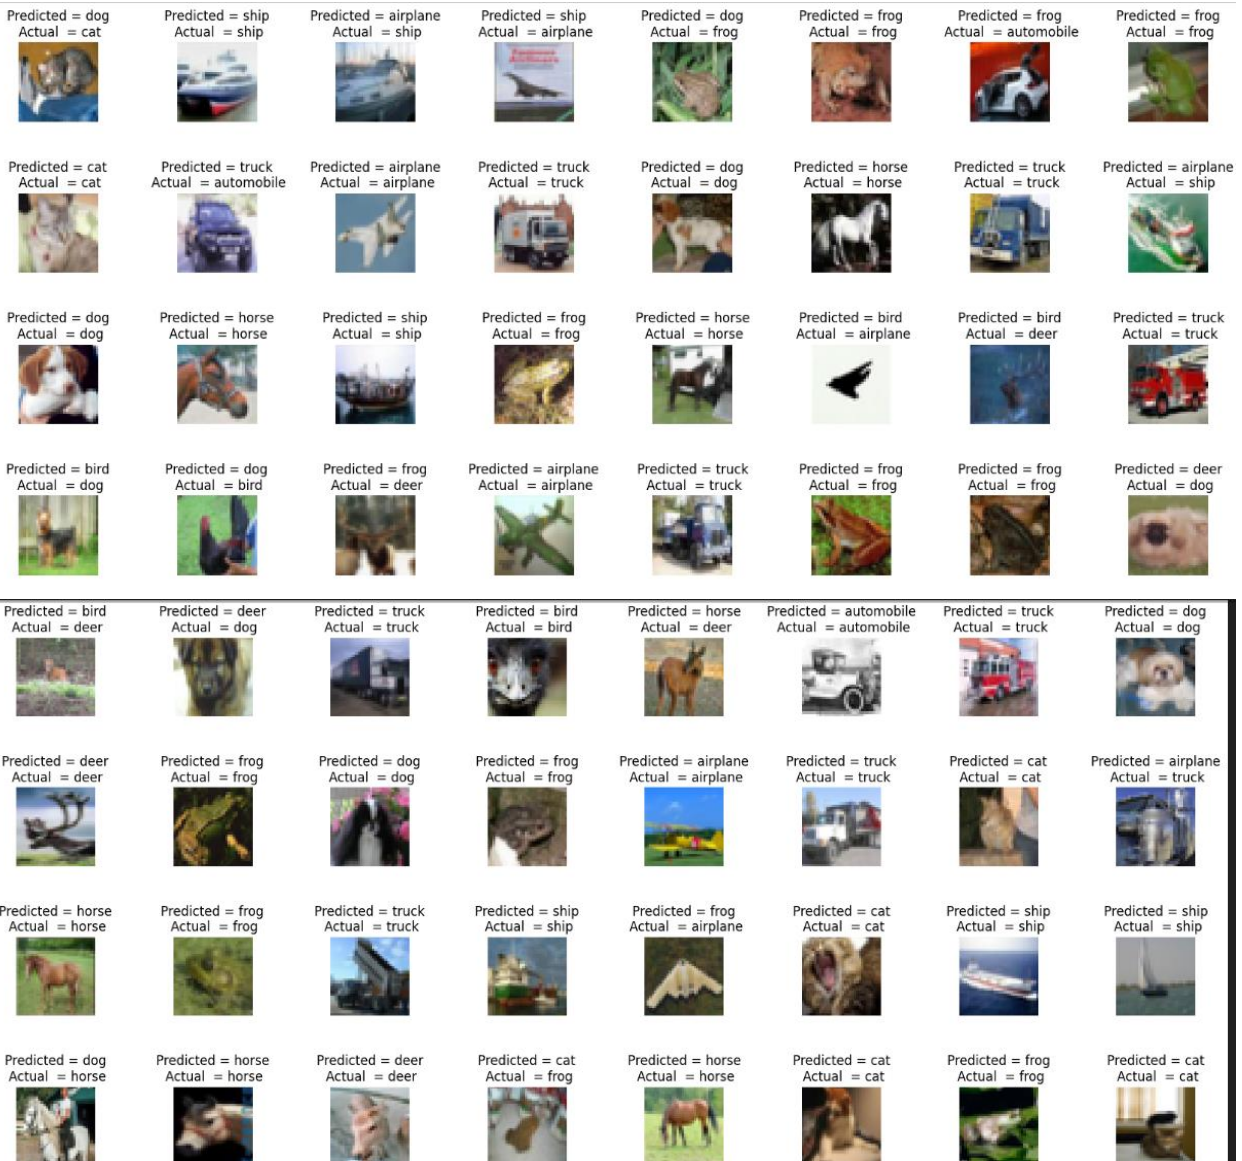
+ Code

+ Markdown

```
L = 8
W = 8
fig, axes = plt.subplots(L, W, figsize = (20,20))
axes = axes.ravel() #

for i in np.arange(0, L * W):
    axes[i].imshow(x_test[i])
    axes[i].set_title("Predicted = {}\n Actual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))
    axes[i].axis('off')

plt.subplots_adjust(wspace=1)
```



3.VGG 16

```
import keras
from keras.models import Sequential
from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D
from keras.datasets import cifar10
from keras import optimizers
from matplotlib import pyplot as plt
```


```
# generate cifar10 data
(x_train,y_train),(x_test,y_test) = cifar10.load_data()
```

```
# config parameters
num_classes = 10
input_shape = x_train.shape[1:4]
optimizer = optimizers.Adam(lr=0.0003)
```

```
# convert label to one-hot
one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)
```

```
# check data
plt.imshow(x_train[1])
print(x_train[1].shape)
```

(32, 32, 3)



```
# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(4096,activation='relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

```
# config optimizer,loss,metrics
model.compile(optimizer=optimizer,loss='categorical_crossentropy',metrics=['accuracy'])
```

```
# check model
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 32, 32, 64)	1792
conv2d_9 (Conv2D)	(None, 32, 32, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_10 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_11 (Conv2D)	(None, 16, 16, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_12 (Conv2D)	(None, 8, 8, 256)	295168

```
conv2d_13 (Conv2D)          (None, 8, 8, 256)          590080
conv2d_14 (Conv2D)          (None, 8, 8, 256)          590080
max_pooling2d_5 (MaxPoolin  (None, 4, 4, 256)          0
...
Total params: 33638218 (128.32 MB)
Trainable params: 33638218 (128.32 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
# train
model.fit(x=x_train,y=one_hot_y_train,batch_size=128,epochs=10)
```

```
Epoch 1/10
391/391 [=====] - 36s 68ms/step - loss: 2.3356 - accuracy: 0.0983
Epoch 2/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0959
Epoch 3/10
391/391 [=====] - 23s 58ms/step - loss: 2.3027 - accuracy: 0.0961
Epoch 4/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0975
Epoch 5/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0969
Epoch 6/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0957
Epoch 7/10
391/391 [=====] - 23s 58ms/step - loss: 2.3027 - accuracy: 0.0967
Epoch 8/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0959
Epoch 9/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0982
Epoch 10/10
391/391 [=====] - 23s 59ms/step - loss: 2.3027 - accuracy: 0.0981
```


```
# evaluate
print(model.metrics_names)
model.evaluate(x=x_test,y=one_hot_y_test,batch_size=512)
```

```
['loss', 'accuracy']
20/20 [=====] - 7s 162ms/step - loss: 2.3026 - accuracy: 0.1000
[2.3025965690612793, 0.10000000149011612]
```

```
# predict
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i,_ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:",predict)
print("expected class:",expect)
```

```
1/1 [=====] - 1s 640ms/step
predict class: 9
expected class: 5
```



```
# save model  
model.save("keras-VGG16-cifar10.h5")
```

4.VGG19Model

```
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import tensorflow as tf  
import keras  
  
from keras.datasets import cifar10  
from tensorflow.keras.optimizers import RMSprop  
from keras.preprocessing import image  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization  
  
%matplotlib inline
```

Extract data and train and test dataset

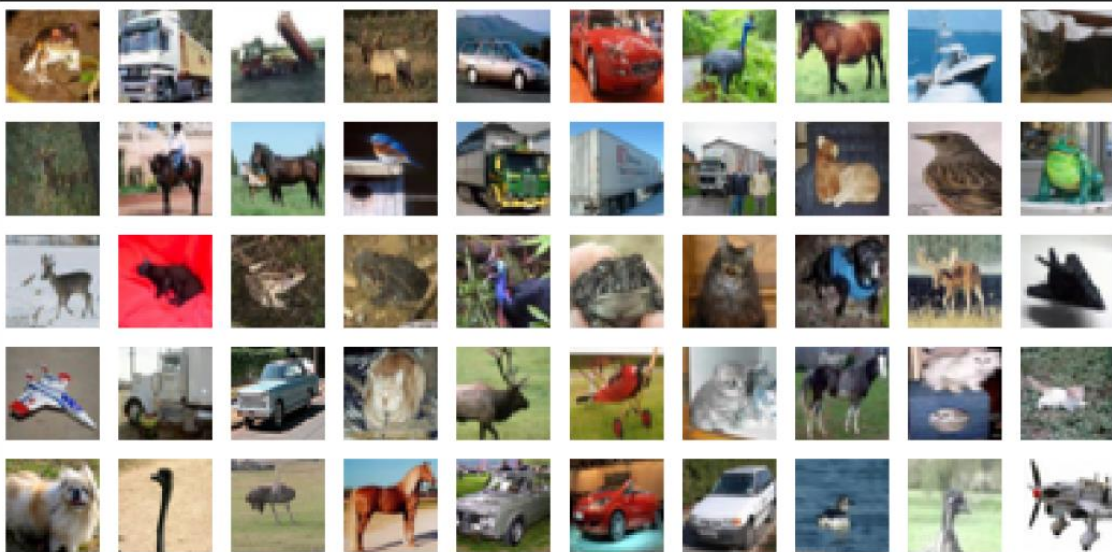
```
cifar100 = tf.keras.datasets.cifar100
(X_train,Y_train) , (X_test,Y_test) = cifar100.load_data()
```


```
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
# X_train = X_train.astype('float32')
# X_test = X_test.astype('float32')
# X_train = X_train / 255.0
# X_test = X_test / 255.0
```

let's look into the dataset images

```
plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,1+i)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```





Training , Validating and Splitting trained and tested data

```
from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train, Y_train, test_size=0.2)
```

```
from keras.utils import to_categorical
y_train = to_categorical(y_train, num_classes = 10)
y_val = to_categorical(y_val, num_classes = 10)
```

```
print(x_train.shape)
print(y_train.shape)
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)
```

```
(40000, 32, 32, 3)
(40000, 10)
(10000, 32, 32, 3)
(10000, 10)
(10000, 32, 32, 3)
(10000, 1)
```

```

train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)

val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)

```

```

from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
                                             factor=0.5,
                                             min_lr=0.00001)

```

We have used only 16 layers out of 19 layers in the CNN

```

vgg_model = tf.keras.applications.VGG19(
    include_top=False,
    weights="imagenet",
    input_shape=(32,32,3),
)

vgg_model.summary()

```

Model: "vgg19"

Layer (type)	Output Shape	Param #

input_3 (InputLayer)	(None, 32, 32, 3)	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
...		
Total params: 20024384 (76.39 MB)		
Trainable params: 20024384 (76.39 MB)		
Non-trainable params: 0 (0.00 Byte)		

```

model = tf.keras.Sequential()
model.add(vgg_model)
model.add(Flatten())
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(1024, activation = 'relu'))
model.add(Dense(256, activation = 'sigmoid'))
model.add(Dense(10, activation = 'softmax'))

model.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten_3 (Flatten)	(None, 512)	0
dense_10 (Dense)	(None, 1024)	525312
dense_11 (Dense)	(None, 1024)	1049600
dense_12 (Dense)	(None, 256)	262400
dense_13 (Dense)	(None, 10)	2570

```

=====
Total params: 21864266 (83.41 MB)
Trainable params: 21864266 (83.41 MB)
Non-trainable params: 0 (0.00 Byte)

```

```

optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer= optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

```

```

history = model.fit(
    train_datagen.flow(x_train, y_train, batch_size = 128),
    validation_data = val_datagen.flow(x_val, y_val, batch_size = 128),
    epochs = 10,
    verbose = 1,
    callbacks = [learning_rate_reduction]
)

```

```


313/313 [=====] - 38s 112ms/step - loss: 1.1032 - accuracy: 0.6178 - val_loss: 0.7488 - val_accuracy: 0.7486 - lr: 0.0010
Epoch 2/10
313/313 [=====] - 34s 107ms/step - loss: 0.7297 - accuracy: 0.7541 - val_loss: 0.6326 - val_accuracy: 0.7842 - lr: 0.0010
Epoch 3/10
313/313 [=====] - 34s 107ms/step - loss: 0.6273 - accuracy: 0.7884 - val_loss: 0.6150 - val_accuracy: 0.7950 - lr: 0.0010
Epoch 4/10
313/313 [=====] - 34s 107ms/step - loss: 0.5704 - accuracy: 0.8069 - val_loss: 0.5969 - val_accuracy: 0.7970 - lr: 0.0010
Epoch 5/10
313/313 [=====] - 34s 108ms/step - loss: 0.5160 - accuracy: 0.8231 - val_loss: 0.5143 - val_accuracy: 0.8258 - lr: 0.0010
Epoch 6/10
313/313 [=====] - 34s 109ms/step - loss: 0.4794 - accuracy: 0.8351 - val_loss: 0.5276 - val_accuracy: 0.8218 - lr: 0.0010
Epoch 7/10
313/313 [=====] - 35s 112ms/step - loss: 0.4492 - accuracy: 0.8448 - val_loss: 0.5124 - val_accuracy: 0.8297 - lr: 0.0010
Epoch 8/10
313/313 [=====] - 33s 106ms/step - loss: 0.4173 - accuracy: 0.8559 - val_loss: 0.4892 - val_accuracy: 0.8359 - lr: 0.0010
Epoch 9/10
313/313 [=====] - 33s 104ms/step - loss: 0.3871 - accuracy: 0.8678 - val_loss: 0.4836 - val_accuracy: 0.8373 - lr: 0.0010
Epoch 10/10
313/313 [=====] - 33s 106ms/step - loss: 0.3677 - accuracy: 0.8744 - val_loss: 0.4372 - val_accuracy: 0.8516 - lr: 0.0010

```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Acuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()
```

<matplotlib.legend.Legend at 0x7f9ed4129930>




```

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure()
plt.plot(loss,color = 'green',label = 'Training Loss')
plt.plot(val_loss,color = 'red',label = 'Validation Loss')
plt.legend()

```

<matplotlib.legend.Legend at 0x7f9ed55d68f0>



```

x_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = np.argmax(model.predict(x_test), axis=-1)
y_pred[:10]

```

313/313 [=====] - 3s 9ms/step

```
array([3, 8, 8, 0, 6, 6, 1, 6, 3, 1])
```

```

cm = confusion_matrix(Y_test, y_pred)
cm

```

```

array([[892,  4, 19,  5, 10,  2,  2, 12, 36, 18],
       [12, 885,  1,  1,  0,  2,  3,  2, 11, 83],
       [35,  1, 826, 19, 48, 19, 35,  9,  5,  3],
       [ 5,  3, 36, 700, 32, 145, 39, 19,  6, 15],
       [12,  1, 37, 27, 834, 15, 29, 40,  4,  1],
       [ 3,  1, 28, 105, 22, 793, 11, 34,  1,  2],
       [10,  2, 28, 38, 15, 11, 884,  1,  7,  4],
       [ 5,  0, 13, 21, 25, 30,  1, 898,  2,  5],
       [29, 15,  5,  1,  3,  1,  0,  0, 923, 23],
       [15, 16,  3, 15,  1,  1,  1,  7, 12, 929]])

```

```
from sklearn.metrics import confusion_matrix, accuracy_score
print('Testing Accuracy : ', accuracy_score(Y_test, y_pred))
```

Testing Accuracy : 0.8564

```
import itertools
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Greens):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=30)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

    #print(cm)

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")
```

```
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

